




INVESTOR OVERVIEW · SEPTEMBER 2026

LunarFlow

In Sync With Your Cycle.

A menstrual tracker that **asks more, claims less**, and can actually explain what it sees — built, tested, and encrypted on the device before a single slide was written.



Drawn to scale: real day counts for a 28-day cycle, not equal thirds.

Ask deeply. Predict honestly. Explain personally.

The category leader is worth **\$1B** and has **77 million** monthly users. It also settled with the FTC for sharing health data it promised to keep, and in September 2025 paid into a **\$56M** class-action settlement alongside Google. Meanwhile the published literature finds that period apps get ovulation exactly right **8% of the time** — and that not one of them tells you so.

The demand is proven. The trust and the accuracy are not. LunarFlow is built on the three commitments the incumbents broke.

| Every figure in this deck is sourced on the final slide. Numbers we do not have are shown as blanks, not estimates.

A \$2.5B market compounding at 20% a year

\$2.49B

Menstrual health apps market,
2026

\$13.1B

Projected 2035

20.3%

CAGR to 2035

72.9%

Of 2024 category revenue came
from **Android**

That last number is the one that matters operationally. LunarFlow is **Android-first by construction** — a \$25 one-time Play fee against Apple's \$99/year — which means the cheap platform and the majority-revenue platform are the same platform.

Forecasts vary by house: SNS Insider puts 2032 at \$7.52B, Towards Healthcare puts 2035 at \$13.11B. Both agree on the growth shape.

One competitor already proved the ceiling

\$1B

Flo valuation, 2024 — first purely digital consumer women's health app to reach unicorn status

77M

Monthly active users

300M+

Registered users; 48M installs in 2025 alone

Flo raised **\$200M from General Atlantic** on those numbers. Nobody needs convincing that people will install a cycle tracker and hand it intimate data.

A note on rigour: reported 2025 revenue disagrees by source — Latka says \$157.6M ARR, Business of Apps says \$275M. We cite both rather than picking the flattering one.

The category spent its trust

2021 · FEDERAL TRADE COMMISSION

Flo promised health data would stay private and be used only to run the app. The FTC alleged it disclosed that data — from millions of users — to **Facebook Analytics, Google Analytics, Fabric, AppsFlyer and Flurry**. The settlement forced an independent privacy review and consent before sharing.

SEPTEMBER 2025 · CLASS ACTION

Google agreed to pay **\$48M** and Flo **\$8M — \$56M combined** — to resolve claims over period-tracking data shared with Google and Meta.

Four years after the FTC order, the same category, the same failure mode.

The biggest brand in the category has taught 300 million people that a period app cannot be trusted with a period. **That is an acquisition asset for whoever arrives second and can prove otherwise.**

And the predictions are mostly wrong

8%

Of ovulation predictions exactly correct, across 10 popular apps

67%

Ran **2–9 days early** for women whose cycle length is non-standard

22%

Of period-tracking apps contain inaccurate information or tools

0

Apps that warned the prediction may be wrong and is only a guide

The last column is the opening. Being wrong is forgivable — cycles are noisy and every app is running a calendar estimate. **Being wrong silently is not.** An entire category has chosen confident-and-wrong over honest-and-useful, and the research says users pay for it in anxiety.

"A false sense of personalization"

A qualitative study of women's real experiences catalogues the harms in the users' own terms: inaccurate predictions, alarmist notifications, a false sense of personalization, and unethical marketing.

When the app's prediction is wrong, users doubt their own body rather than the app.

That is the sentence the product is built against. An app that speaks about a body it knows nothing about is not personalized — it is *guessing in a confident voice*. The fix is not a better guess. It is knowing more, and admitting the rest.

Three counter-moves, all already shipped

01 · ASK DEEPLY

An **11-page intake** — the deepest cold start in the category — plus daily logging across flow, symptoms, mood, pain, temperature, ovulation tests, discharge, sleep, energy, stress and medication adherence.

Day one is not a blank calendar.

02 · PREDICT HONESTLY

No fertility percentage, ever. Predictions gate on confidence and *disappear* below it rather than degrading. A persistent disclaimer sits on every fertility surface — the exact thing the literature found nobody does.

03 · EXPLAIN PERSONALLY

The user's **full 90-day record travels with their question** to a language model — under strict rules about what may and may not be sent, each one pinned by a test.

How LunarFlow draws a prediction: a soft-edged band, never a hard date. The falloff is the uncertainty — the one visual no competitor offers.

Eleven pages, every answer required

- 1. Welcome
- 2. Your data, on your terms
- 3. When did your last period start?
- 4. How long is your cycle, usually?
- 5. When were you born?
- 6. Height, weight, age at first period
- 7. Contraception method
- 8. Sex frequency — typical, and today
- 9. Sexual history, libido
- 10. Solo intimacy, satisfaction
- 11. What are you using LunarFlow for?

"REQUIRED" DOES NOT MEAN TRAPPED

Every question that might not apply carries a **first-class escape answer** — *None · None of these · Not today · Prefer not to say* — and each is stored as a **real key**, not as an empty field.

That distinction is load-bearing. **An absent answer and a negative answer are different clinical facts.** "I have never had pain during sex" is information. "I was never asked" is not. The app never confuses the two — and the AI context builder omits the second rather than inventing it.

Required intake landed 2026-09-14. Before that a skip stored NULL; the wizard now refuses to advance and says why.

Nothing is asked "in case it's useful"

WHAT WE ASK	WHY IT IS REQUIRED	WHAT BREAKS WITHOUT IT	ESCAPE ANSWER
Email + password	Account identity — cross-device restore, and a real identity to erase on request.	No restore after a lost phone; nothing concrete to honour a deletion request against.	A local-only hatch, for users who want no account at all.
Last period start	The anchor date every prediction is computed from.	No forecast at all on day one — an empty calendar until a cycle accrues.	None. The one question the product cannot work without.
Usual cycle length	The interval the first forecast extrapolates on, until real observed cycles replace it.	Month one has no basis for a prediction.	Pre-set to 28 — a real default, not a blank.
Date of birth	Age contextualises every number in the clinician summary, and bounds what counts as a plausible cycle.	The doctor report loses the first line any clinician reads.	Range is 8–100; an implausible date is refused, never clamped .
Height + weight	The only clinical readout computable from profile facts alone — and body composition is a documented correlate of cycle irregularity.	That readout drops out of both the doctor summary and the AI's context.	A wrong value is refused with a reason , not quietly accepted.
Age at first period	A standing gynaecological fact the app can never derive from tracked data.	The summary permanently loses a field every gynaecologist asks about.	"I'm not sure" — stored as unanswered, never guessed.
Contraception method	The highest-value question in the wizard. Ovulation-suppressing methods are held in a named set, and picking one <i>caps prediction confidence to low</i> — suppressing the fertile window app-wide.	The app would show a fertile window to somebody who does not ovulate — asserting something false about their body .	"None" is a genuine answer, not a skip.

Copper IUD and sterilisation are deliberately outside the suppressing set: ovulation continues, and blanking the window would delete a real signal.

The intimate questions, and their defence

WHAT WE ASK	WHY IT IS REQUIRED	WHAT BREAKS WITHOUT IT	ESCAPE ANSWER
Sex frequency — typical and today	Two facts at two time scales. The baseline is the denominator; today's answer is an event that becomes a real log entry.	Any correlation between activity and symptoms has no baseline to be measured against.	"Never" and "Not today" are both real answers.
Sexual history — pain, bleeding after sex, dryness	Gynaecological red flags clinicians actively screen for. They reuse the daily-chip keys, so baseline and daily log can never drift into meaning different things.	The pattern detectors and the doctor summary lose their screening inputs.	"None of these" — an explicit key meaning <i>asked and answered no</i> .
Libido — baseline and today	A hormonal-cycle correlate. The baseline is what gives the daily chip a reference point.	A daily libido chip with no baseline is uninterpretable — high relative to what?	A three-point scale that is always answerable.
Masturbation frequency	Completes the sexual-activity picture the cycle correlations draw on, without requiring a partner to be involved.	Activity data becomes partner-only, under-counting a large share of users.	"Never" is a real answer.
Goal — track, conceive, perimenopause	Changes what the app <i>shows</i> , not just how it looks: conceive promotes fertility above period; perimenopause caps confidence, self-suppressing every fertility estimate.	Perimenopausal users would receive fertility estimates the app has no business making.	Defaults to "Track my cycle".

Two further questions are asked at signup and are *not* on this table, because they cannot be defended on these terms. They are the subject of the next slide.

WHY ASK ANY OF THIS

Three of the five are **screening questions** a clinician asks in a first appointment — pain during sex, bleeding after sex, dryness. An app that never asks them is not being discreet; it is discarding the signals most likely to matter.

WHY ASK IT AT SIGNUP

A daily chip records *today*. It can never establish what is *typical* — and without that, one day's answer has nothing to be unusual against. The baseline is the question a first run can answer and a daily log cannot.

WHY IT IS ANSWERABLE ANYWAY

Every escape answer is a **stored key, not an empty field**. "None of these" means *asked and answered no* — information a clinician can use. A skip would have meant nothing at all.



Two questions we cannot justify

The wizard asks two more required questions — **preferred methods of solo intimacy, and time to satisfaction**. Neither is defensible on clinical grounds, and we know that because *our own source code says so, in a comment written before either was added*:

"Preferred methods and time to orgasm are deliberately absent: they carry no gynaecological signal while being among the most sensitive things this app could store."

WHAT HAPPENED

The judgement was recorded as **correct on both counts**, then overridden by the owner on 2026-09-14 — and the reversal was **documented in the file** so that no future reader mistakes it for an oversight and quietly undoes it.

WHAT THEY COST

Nothing consumes them — **no prediction, report, detector or insight**. They sit at the deepest point of the funnel, on the most sensitive page, producing no product value. Both carry "Prefer not to say", stored as a real key.

WHAT WE DO ABOUT IT

They are **first on the roadmap to cut**, and cutting them is a one-file change. Under GDPR's data-minimisation test this is the weakest ground the intake stands on, and we would rather you heard it here.



What the depth is deliberately *not* used for

Collecting more is only defensible if you also claim less. Every one of these is a shipped constraint, not an intention:

- **No fertility percentage, ever.** A precise figure — "3.2% chance today" — was proposed and rejected. It reads to a user as a *safe day*, and it collides with Apple's medical-claims rule and Play's health-misinformation policy. The app ships a qualitative band instead, or nothing.
- **Predictions vanish rather than degrade.** Below a confidence threshold the ovulation marker, the fertility band and even the legend entry are all suppressed. An estimate the app does not trust is not shown in a weaker font — it is not shown.
- **Perimenopause caps confidence to low** — one lever that suppresses every fertility surface at once — and pairs it with copy that says the quiet part: *"You can still become pregnant — cycle timing is not a reliable guide."* A hidden window must never read as a safe one.
- **A disclaimer sits on every fertility surface, permanently.** Not on first run, not dismissible. The research finding was that *none* of the popular apps do this.
- **Sexual-activity data is namespace-excluded from the doctor report by default.** The report is a file the user hands to another human; sensitive keys are held out of it structurally, not by remembering to filter.
- **Automated test-strip reading was vetoed.** Interpreting an ovulation strip is a medical-device function. Users may photograph and log one; the app will not tell them what it means.

What reaches the model — and what never does

The context builder is a **pure function**: rows in, text out. It cannot reach the database or the network, and each rule below is pinned by a test that fails if the rule is broken.

RULES ON WHAT IS SENT

- **90 days.** About three cycles — the minimum for a before-and-after-period comparison to contain more than one instance.
- **Unanswered fields are omitted entirely** — never sent as `null` or `"unknown"`, because an absent answer and a negative answer are different clinical facts.
- **Unknown option keys are dropped, never printed raw.** A renamed catalogue entry cannot leak a database slug into a prompt.
- **Zero means unset** for every numeric metric, and a zero is never serialised as a measurement.

RULES ON WHAT IS SAID BACK

- **The model is instructed it is not a clinician:** never diagnose, never name a condition, never estimate severity, never advise treatment — and to refuse and redirect if asked.
- **A caveat renders under every answer, without exception** — because *"a caveat that appears only sometimes teaches the user that its absence means the answer is trustworthy."*
- **Opt-in, off by default,** behind an explicit consent sheet, versioned so the consent can be re-asked when the terms change.
- **Capped** at 20 messages per day per user — which also caps the marginal cost.

Refusal behaviour is re-verified by hand against escalating adversarial prompts: a unit test proves only that the instruction's clauses exist, not that the model still obeys them under pressure.



Why intake depth *is* the AI product

WHAT A TYPICAL TRACKER CAN TELL A MODEL

Age: —
Cycle: ~28 days
Last period: 2026-09-02

Any model given this can only produce generic health writing. That is what "a false sense of personalization" looks like from the inside.

WHAT LUNARFLOW CAN

Age · height · weight · body-composition readout
Age at first period
Contraception method + start date
Clinician-given diagnoses · breastfeeding status
— and 90 days of —
flow · symptoms · mood · pain 0–10 · temperature · ovulation
tests · cervical mucus · sleep · energy · stress · hydration ·
medication adherence

Same model. Same price. The second one stops guessing. The intake solves the cold start, the daily log compounds it, and none of it is portable to a competitor — a user would have to re-answer eleven pages and re-live ninety days to rebuild it somewhere else.

This is built, not designed

143

Dart source files

150

Test files

55

Service modules

v10

Database schema, migrated in place

Shipped and working

- Daily logging + calendar
- Predictions + multi-month forecast
- Insights + plain-language narration
- Doctor-facing PDF report
- Reminders, incl. custom
- Medications + adherence
- Pregnancy mode (loss-safe)
- Perimenopause mode
- Symptothermal temperature + tests
- Health Connect import
- Home-screen widget
- App lock + biometrics
- Diary, media, AI descriptions
- Accounts + cross-device sync
- Backup, restore, account deletion
- Ads + one-time Premium purchase

DEVICE-VERIFIED, NOT ASSUMED

The on-device database is **encrypted at rest** with a 256-bit key held in the OS keystore — confirmed on physical hardware, because the cipher only compiles in a device build and *a green test suite proves nothing here*.

Evidence: the database file begins `3a40 3cd2...`, not the plaintext `SQLite format 3` magic — with the older unencrypted file retained as a positive control proving the check can tell the difference.

The privacy decision is also the cost decision

LOCAL-FIRST BY CONSTRUCTION

- The on-device database is the **single source of truth**. The interface never reads from the cloud.
- Sync runs **alongside** the read/write path, never in front of it — so the app is fully usable offline, permanently.
- Every prediction, insight and report is **computed on the phone**. There is no inference server to run, scale, or pay for.

WHAT A USER ACTUALLY COSTS

- No analytics stack, no recommendation service, no model hosting.
- Marginal cost $\approx$ cloud reads/writes + **AI calls capped at 20 messages per user per day**.
- Fixed cost is a **\$25 one-time** Play registration. Apple's \$99/year is deferred until iOS is worth it.
- Revenue: non-personalized ads — **banned from the logging and insights screens** — plus a one-time purchase that removes them.

Ads are barred from the two screens where a user is most vulnerable and most engaged. That costs inventory. **It is the price of the sentence on slide five being true about us.**

What is wrong with this today

- 1 **The intake is a funnel risk.** Eleven required pages including sexual history is the deepest cold start in the category, and two of those questions are documented *in our own source* as carrying no gynaecological signal. Mitigation: they are the first candidates to cut, and cutting them is a single-file change.
- 2 **Cloud sync is not end-to-end encrypted.** Logs and settings are stored in plaintext under the user's account; the operator can read them. The privacy policy says so plainly. For a trust-led brand this belongs on the roadmap, not in the footnotes.
- 3 **The AI key currently ships inside the app.** It is recoverable from the installed binary with one command. A server-side proxy is a documented *release blocker*, not an accepted state — but it is not built yet.
- 4 **Nothing has launched.** Ads still run on the vendor's test identifiers; the store listing needs a 12-tester, 14-day closed test before a new personal account can publish. Revenue today is zero and the deck does not pretend otherwise.
- 5 **This is special-category data.** Sex life and health data carry the highest bar under GDPR — explicit consent plus data minimisation. Both principles point at the same place as risk 1: trim the intake to what is clinically defensible.

| A deck that omits this list is not lower-risk. It is the same risk, discovered later, by you.

Ordered by what unblocks revenue

NEXT · TO LAUNCH

- Trim or soften the two unjustified intake questions
- Server-side proxy for the AI key
- Real ad and in-app-purchase identifiers
- Closed test, store listing, data-safety forms

THEN · TO DEFEND THE POSITION

- End-to-end encryption for sync
- Translations — the infrastructure is built and deliberately English-only, because machine-translating a phrase like *"not a contraceptive method"* is a safety risk, not a chore
- iOS, once the \$99/year is earned

LATER · TO EXTEND THE MOAT

- Deeper symptothermal confidence integration
- Wearables beyond temperature
- Longitudinal pattern nudges across years, not cycles

Note the order. **The first thing on the list is removing data we already collect** — which is the whole argument of this deck, applied to ourselves.



What we are raising, and what it buys

TERMS — TO BE COMPLETED

- Raising — — — — —
- At — — — — — pre-money
- Instrument — — — — —
- Runway — — — — — months

TRACTION — TO BE COMPLETED

- Installs — — — — —
- D30 retention — — — — —
- Intake completion rate
— — — — —
- Premium conversion
— — — — —

WHAT IT FUNDS

- Launch: proxy, store compliance, closed test
- End-to-end encryption for sync
- Human translation of safety-critical copy
- iOS

These blanks are deliberate. Every other number in this deck is sourced and checkable; inventing these four would undo that.

Every figure, checkable

Market

- Towards Healthcare — *Menstrual Health Apps Market Sizing* (\$2.49B 2026 → \$13.11B 2035, 20.28% CAGR; Android 72.9% of 2024 revenue). towardshealthcare.com/insights/menstrual-health-apps-market-sizing
- SNS Insider via GlobeNewswire, 6 Nov 2025 — \$7.52B by 2032.
- FemTech World — *Menstrual health apps market tipped to reach US\$13bn*.

Competitor scale

- General Atlantic press release — *Flo Health Secures More than \$200M* (\$1B valuation, first purely digital consumer women's health app to reach unicorn status).
- Business of Apps — *Flo Revenue and Usage Statistics (2026)* (77M MAU; 300M+ registered; 48M installs in 2025; \$275M revenue).
- Latka — Flo Health (\$157.6M ARR). Cited alongside the above because the two disagree.

Trust failures

- Federal Trade Commission, Jan 2021 — *Developer of Popular Women's Fertility-Tracking App Settles FTC Allegations* (disclosure to Facebook Analytics, Google Analytics, Fabric, AppsFlyer, Flurry). ftc.gov
- Global Policy Watch, Oct 2025 — *Flo Health, Google Settle Class Action Privacy Lawsuit for \$56 Million* (Google \$48M, Flo \$8M).

Accuracy & harm

- Worsfold, Marriott, Johnson & Harper — *Period tracker applications: what menstrual cycle information are they giving women?*, Women's Health (2021) — 8% of ovulation predictions exactly correct; 67% ran 2–9 days early on non-standard cycles.
- *The need for transparency in the promotion of popular period tracker applications* — app-store content analysis (2025): 22% contained inaccurate information; none caveated their own unreliability.
- *The benefits and harms of using a period tracking app: a qualitative study of women's experiences*, ScienceDirect (2026) — "inaccurate predictions, alarmist notifications, a false sense of personalization, and unethical marketing"; users doubting their body before the app.